


# A A prediction is assembled from many models, not produced by one



## B One endpoint is twenty-one fits

The panel above counts the estimators that answer a query. Each was preceded by twenty that never serve a prediction and exist only to measure how the twenty-first behaves on compounds it has not seen.

10 random folds



compounds partitioned at random;  
reports interpolation within known chemistry

10 scaffold folds



whole Bemis-Murcko scaffolds held out;  
reports generalisation to new chemistry

1 deployed fit



refitted on every remaining row, then calibrated; this is what the server runs

Each block is one random forest of 300 trees, so the panel rests on far more fitted estimators than it deploys, and every reported interval is measured on compounds withheld from the fit that produced it.

70

deployed estimators, plus  
8 isotonic calibrators

1,480

cross-validation fits, over  
74 cross-validated estimators

300

trees in every forest,  
leaf size 2

1,036

input columns, identical  
for every endpoint