


A A prediction is assembled from many models, not produced by one



B One endpoint is twenty-one fits

10 random folds



compounds partitioned at random;
reports interpolation within known chemistry

10 scaffold folds



whole Bemis-Murcko scaffolds held out;
reports generalisation to new chemistry

1 deployed fit



refitted on every remaining row, then
calibrated; this is what the server runs

70

deployed estimators, plus
8 isotonic calibrators

1,480

cross-validation fits, over
74 cross-validated estimators

300

trees in every forest,
leaf size 2

1,036

input columns, identical
for every endpoint