A **A prediction is assembled from many models, not produced by one**

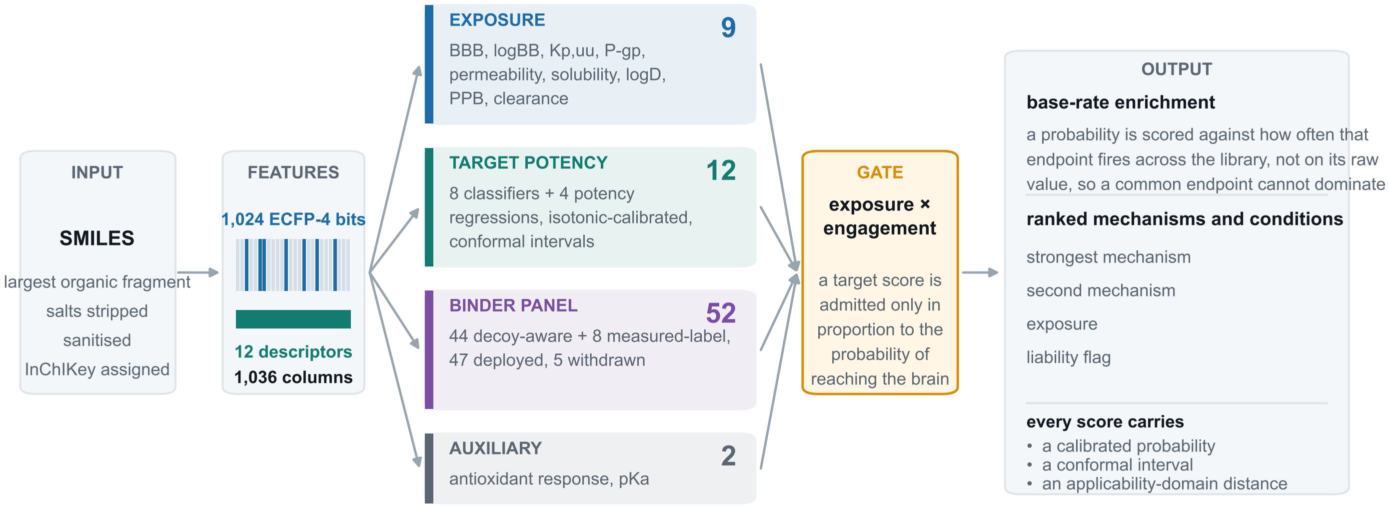

B

One endpoint is twenty-one fits

The panel above counts the estimators that answer a query. Each was preceded by twenty that never serve a prediction and exist only to measure how the twenty-first behaves on compounds it has not seen.

compounds partitioned at random; reports interpolation within known chemistry

whole Bemis-Murcko scaffolds held out; reports generalisation to new chemistry

refitted on every remaining row, then calibrated; this is what the server runs

Each block is one random forest of 300 trees, so the panel rests on far more fitted estimators than it deploys, and every reported interval is measured on compounds withheld from the fit that produced it.

70 deployed estimators, plus 8 isotonic calibrators	**1,480** cross-validation fits, over 74 cross-validated estimators	**300** trees in every forest, leaf size 2	**1,036** input columns, identical for every endpoint