




A calibration works on the core



B and the gate is the worst of them



C conformal coverage, two endpoints short



D where each layer stops

	8 core classifiers	38 binder endpoints	ADME and auxiliary
isotonic calibration	yes	no	no
conformal set	yes	no	no
applicability band	yes	yes	yes
per-endpoint calibration error	yes	yes	no

binder calibration error: mean 0.0762, median 0.0675, range 0.0300 to 0.1780
against 0.0147 for the core, about five times larger, on 38 of the 47 deployed endpoints

The stack is fullest for the eight endpoints a user rarely queries and thinnest for the thirty-eight they query constantly. Extending it to the panel needs no new data and is the largest single improvement available.