


A ten hypotheses, four refuted and two weakened

hypothesis, as posed	verdict	what the refutation cost
H1 the disease score is informative	SUPPORTED	
H2 the curated edge weights add value	REFUTED	the graph is described as structure, not tuned parameters
H3 BBB gating discriminates between diseases	REFUTED	by construction. The gate is a filter, not a discriminator
H4 specificity transfers to novel chemistry	SUPPORTED	
H5 read-across beats a frequency baseline	SUPPORTED	
H6 the disease scores match clinical indications	WEAKENED	
H7 some panel targets cannot rank at all	REFUTED	four endpoints withdrawn for firing on glucose and urea
H8 engaged targets are independent observations	REFUTED	the interface now groups homologues and quotes the correlation
H9 the disease layer discriminates between compounds	SUPPORTED	
H10 the barrier model earns its place over a rule	WEAKENED	


H10 was added after a thesis audit found that the component the architecture is named for had no hypothesis.

B the leakage null: labels permuted



circle = random split, square = scaffold. permuted means 0.4959 and 0.5026, all sixteen within 0.0200 of chance

C H10, the hypothesis the suite lacked



on 241 approved drugs the model has never seen, the strongest null sits 0.0233 behind with an interval reaching +0.0106, so the verdict is WEAKENED

The suite is read-only: no model, dataset or threshold was changed to obtain any number in it, and acting on a finding is a separate commit, so a wording change cannot be mistaken for evidence.